

Pod Name	Node	GPU	Startup (sec)
qwen-qwe-c9c8568d-en3-0-6b-decode-7bbc774556-v95js	pokprod-b93r39s2	NVIDIA-H100-80GB-HBM3	148.0s
qwen-qwe-c9c8568d-en3-0-6b-decode-7bbc774556-gccwx	pokprod-b93r43s2	NVIDIA-H100-80GB-HBM3	180.0s

EPP Configuration (Feature Gates & Scorer Weights)

apiVersion: inference.networking.x-k8s.io/v1alpha1
kind: EndpointPickerConfig
featureGates:
- flowControl
plugins:
- type: queue-scorer
- type: kv-cache-utilization-scorer
- type: prefix-cache-scorer
schedulingProfiles:
- name: default
 plugins:
 - pluginRef: queue-scorer
 weight: 2
 - pluginRef: kv-cache-utilization-scorer
 weight: 2
 - pluginRef: prefix-cache-scorer
 weight: 3


Benchmark Load Generator Configuration

Profile: chatbot_synthetic.yaml
target: http://infra-llmdbench-inference-gateway-istio.wva-bench-test.svc.cluster.local:80
model: Qwen/Qwen3-0.6B
request_type: text_completions
profile: constant
rate: [1,2,4,8]
max_seconds: 600
random_seed: 42
data:
 prompt_tokens_min: 10
 prompt_tokens_max: 8192
 prompt_tokens: 4096
 prompt_tokens_stdev: 2048
 output_tokens_min: 10
 output_tokens_max: 2048
 output_tokens: 1024
 output_tokens_stdev: 512
 samples: 1000


Autoscaling Configuration (HPA & VA)

Variant (VA) [workload-variant-autoscaler-va]: Min Replicas = 1, Max Replicas = 10, Cost Factor = 10.0
HorizontalPodAutoscaler (HPA) [workload-variant-autoscaler-hpa]: Min Replicas = 1, Max Replicas = 10
> ScaleUp: Stabilization Window = 240s | Policies = [{"periodSeconds": 15, "type": "Percent", "value": 100}]
> ScaleDown: Stabilization Window = 240s | Policies = [{"periodSeconds": 15, "type": "Percent", "value": 100}]

KV Cache Usage Percentage Over Time (Namespace: wva-bench-test)



Number of Requests Waiting & Running Over Time (Namespace: wva-bench-test)



Decode Replica Count & EPP Queue Over Time (Namespace: wva-bench-test)

